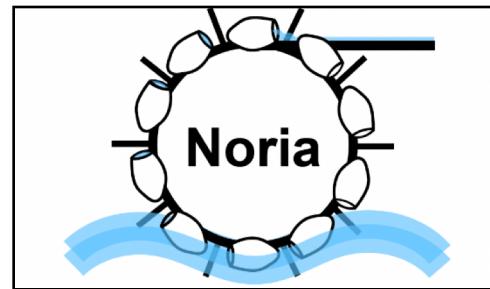


Noria: Dynamic Partially-stafeful data flow for high-performance web applications

Gjengjet et al. (OSDI'18)



Presentation by **Yannis Marketakis**

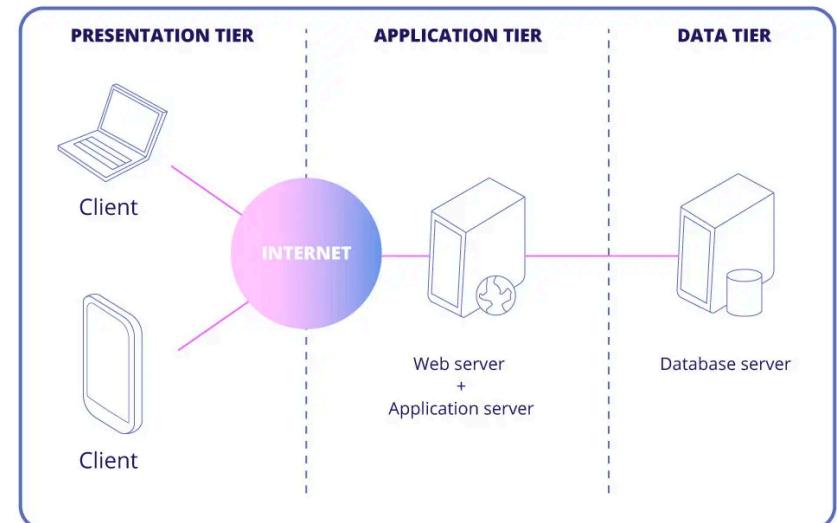
In the context of HY559 (Computer Science Department, UOC)

Outline

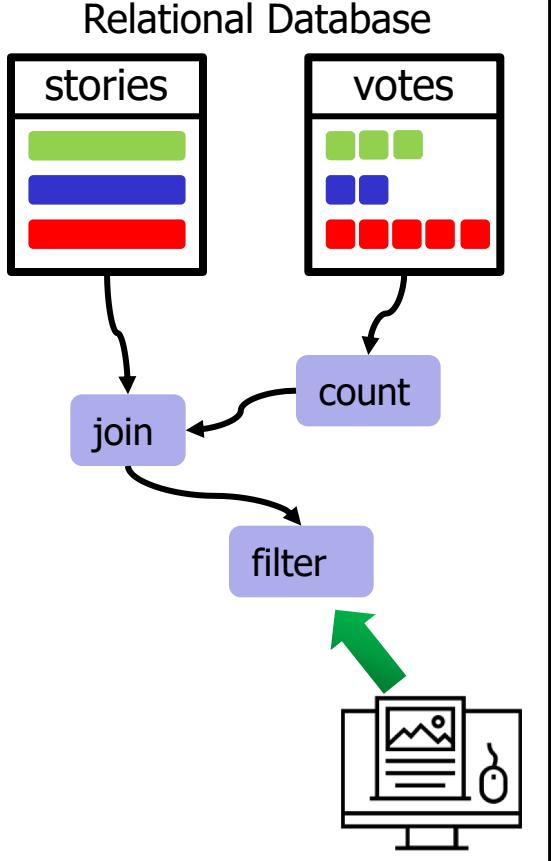
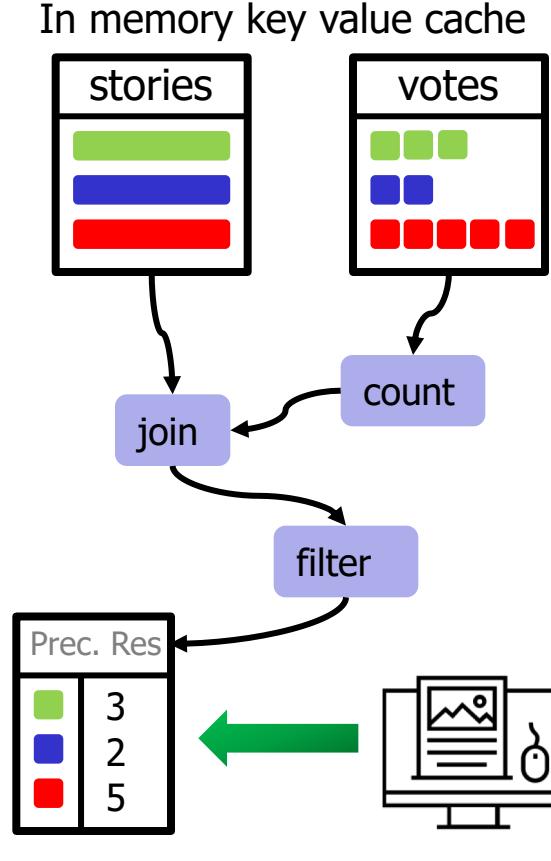
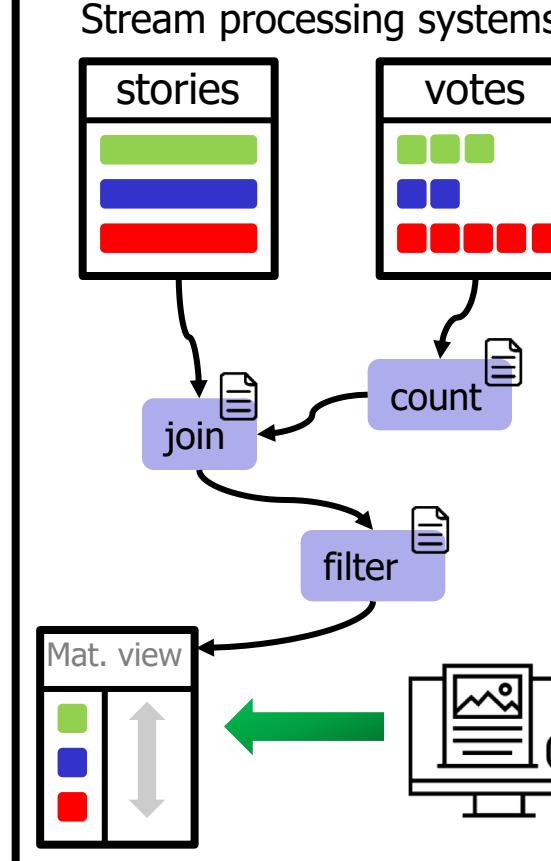
- Motivation - Background [3'] (sl. 3-4)
- Noria [10'] (sl. 5-12)
 - Introduction
 - Partially-stateful data flow
 - Dynamic data flow
- Implementation Details [2'] (sl. 13-14)
- Evaluation [6'] (sl. 15-20)
- Conclusion [1'] (sl. 18)

Motivation - Background

- Web applications must serve users at low latency
- Responses are computed by querying data from backend stores (i.e. relational databases)
- Typically, web applications are read-heavy
 - Lobste.rs: 97% reads
 - HotCRP: 88% reads
 - Facebook: 99.8 % reads
 - Bronson et. al. (2013)



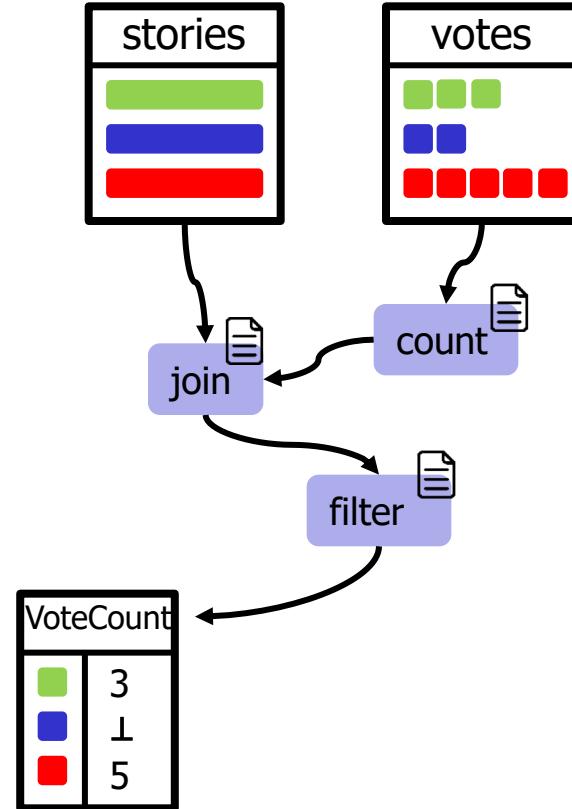
Motivation & Background – cont'd

Relational Database	In memory key value cache	Stream processing systems
 <p>Relational Database</p>	 <p>In memory key value cache</p>	 <p>Stream processing systems</p>
<p>Increased query execution time Repeated work (same queries)</p> <p>:(</p>	<p>Updates → manual modification, or cache invalidation Complex application-side logic</p>	<p>State can grow rapidly New queries requires restarting</p>

Noria

- Noria combines the best of the three alternatives
 - Fast reads (in-memory key-value cache)
 - Efficiency against updates and parallelism (stream-processing systems)
 - Changing queries and view without downtime (relational databases)

- Key Characteristics
 - Bound memory footprint
 - Live addition/updates of new queries/views
 - Without global coordination

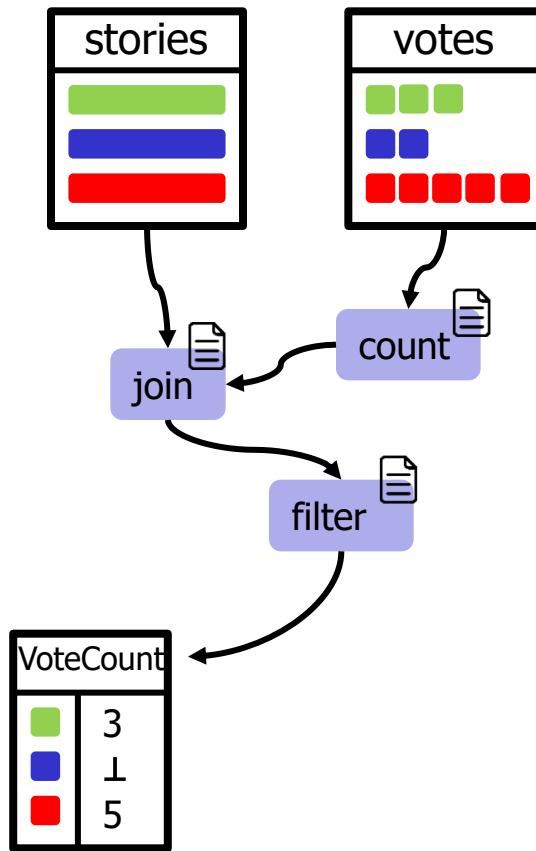


Noria – Challenges

- Limit the size of state and views
 - Partially-stateful data-flow
- Changes to Noria program must adapt the data-flow without downtime
 - Dynamic data-flow

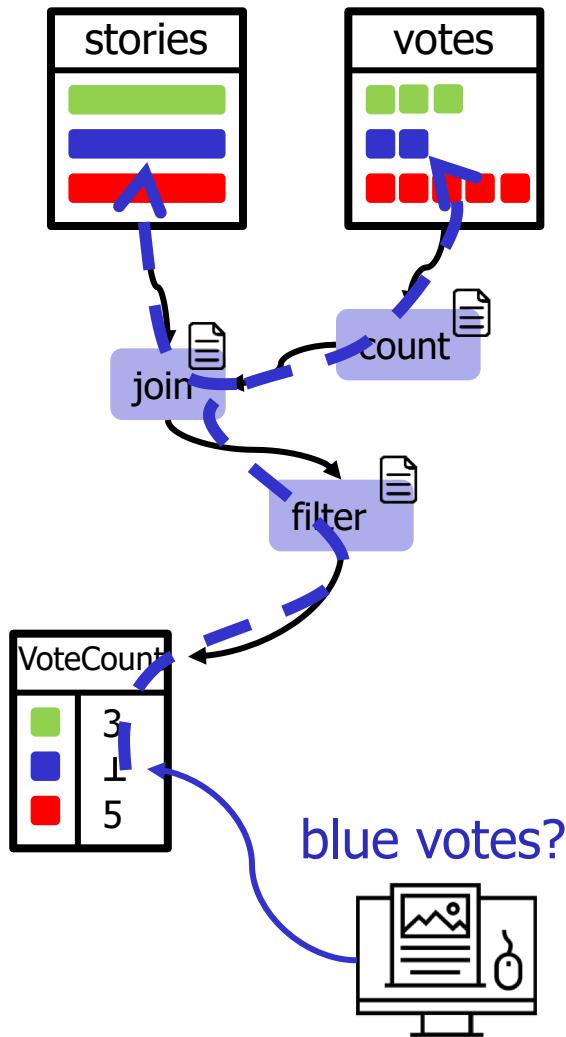
Noria – Partially-stateful data-flow

- Operators maintain only a subset of their state
 - Absent entries (\perp)
- Why absent?
 - They were never filled in (i.e. not requested)
 - They were evicted by the system (i.e. rarely used)
- Benefits
 - It reduces the memory footprint efficiently
 - It does not rely on a “window cache”



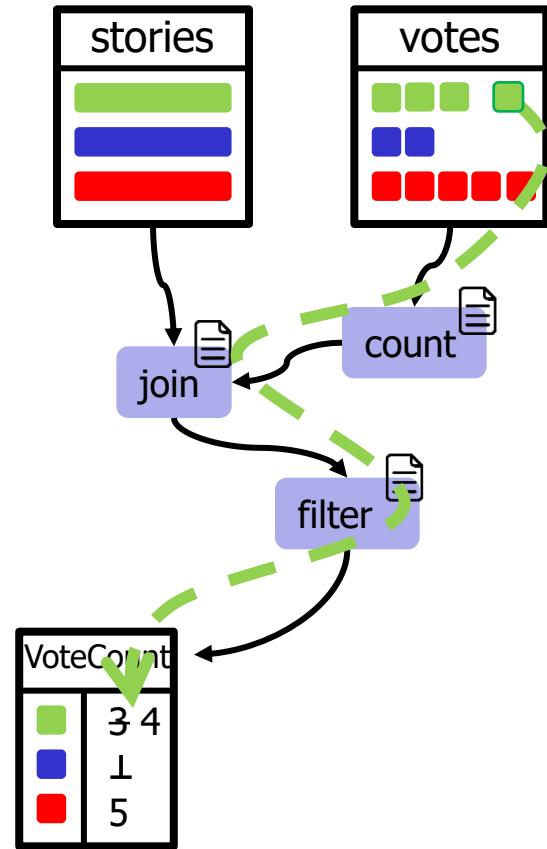
Noria – Partially-stateful data-flow

- Read from evicted state?
 - Upquery the data flow stream
- Fill the evicted state
 - In the operators
 - In the external views



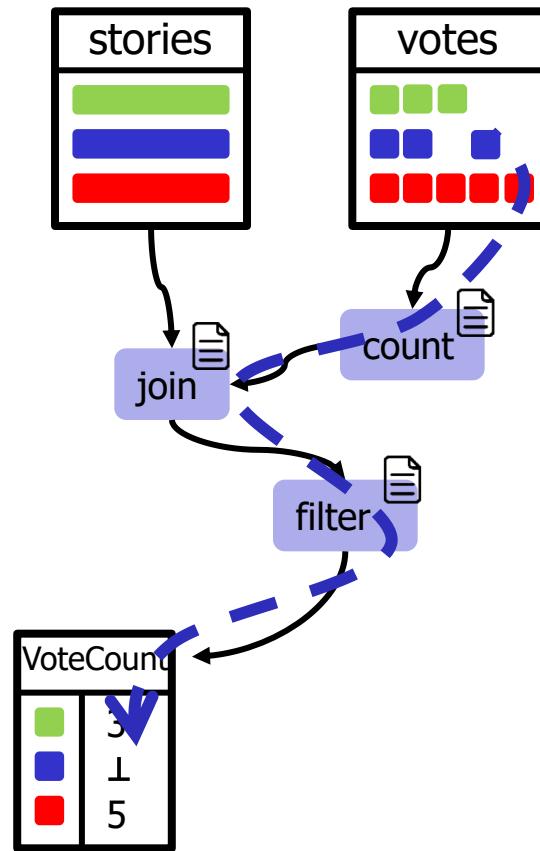
Noria – Partially-stateful data-flow

- Writes
 - Applied to a base table and injected in the data flow as an update
 - Eventually updates read and modify external views
- Example
 - A new vote for the green story
 - Stored in the base table (votes)
 - State updated
 - Operators of the data flow
 - External view



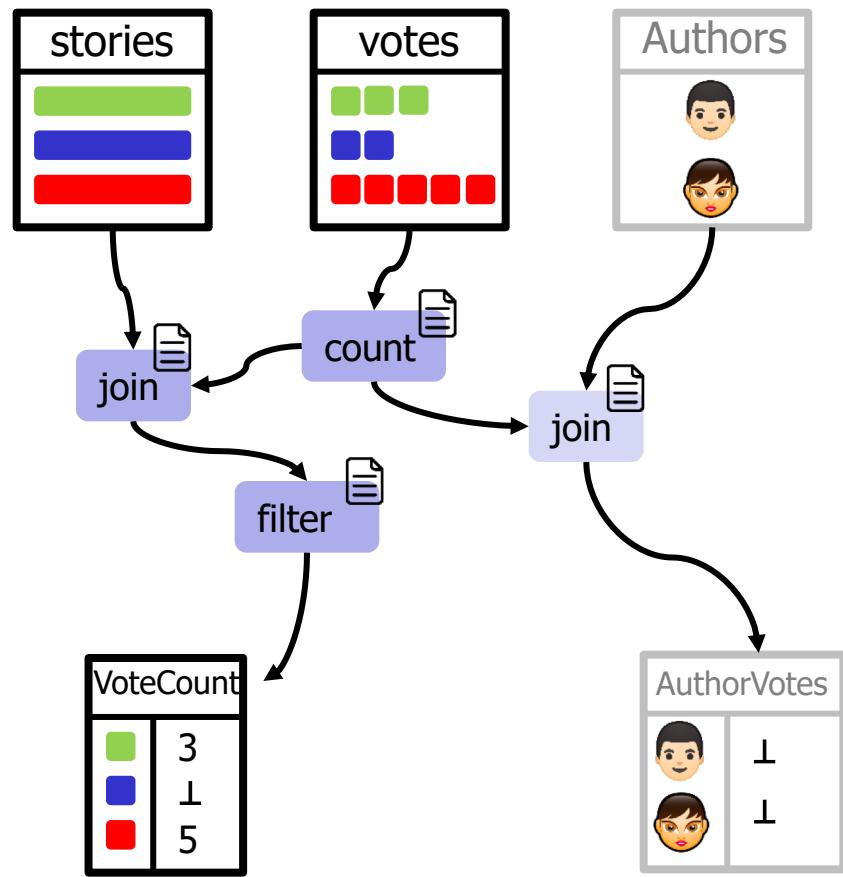
Noria – Partially-stateful data-flow

- If a request for updating an evicted state arrives (\perp) → drop it
- Benefits
 - There is no need to update missing entries
 - They will be updated as soon as they are requested (via an upquery)
- Example
 - A new vote for the green story
 - Stored in the base table (votes)



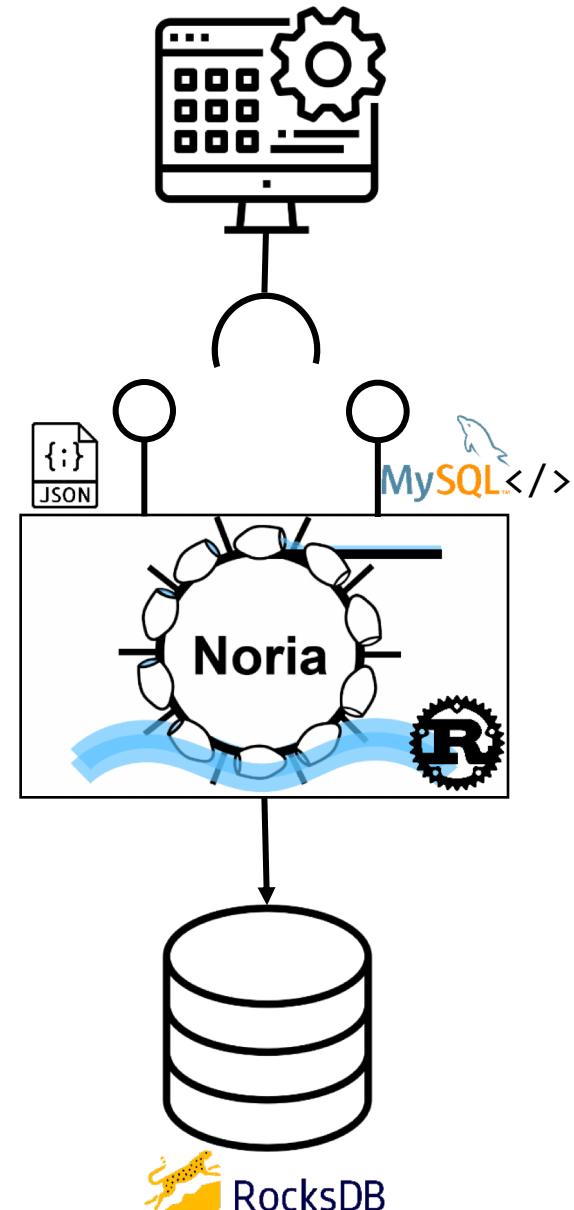
Noria – Dynamic data-flow

- 1st step: transition planning
 - The app provides added and removed expressions
 - Noria computes the required changes (this allows re-use of operators)
- 2nd step: data flow transition
 - Adds new operators
 - Originally their state is empty
 - Remove obsolete operators
- Challenges
 - Avoid state duplication
 - Continuous supports of reads and writes

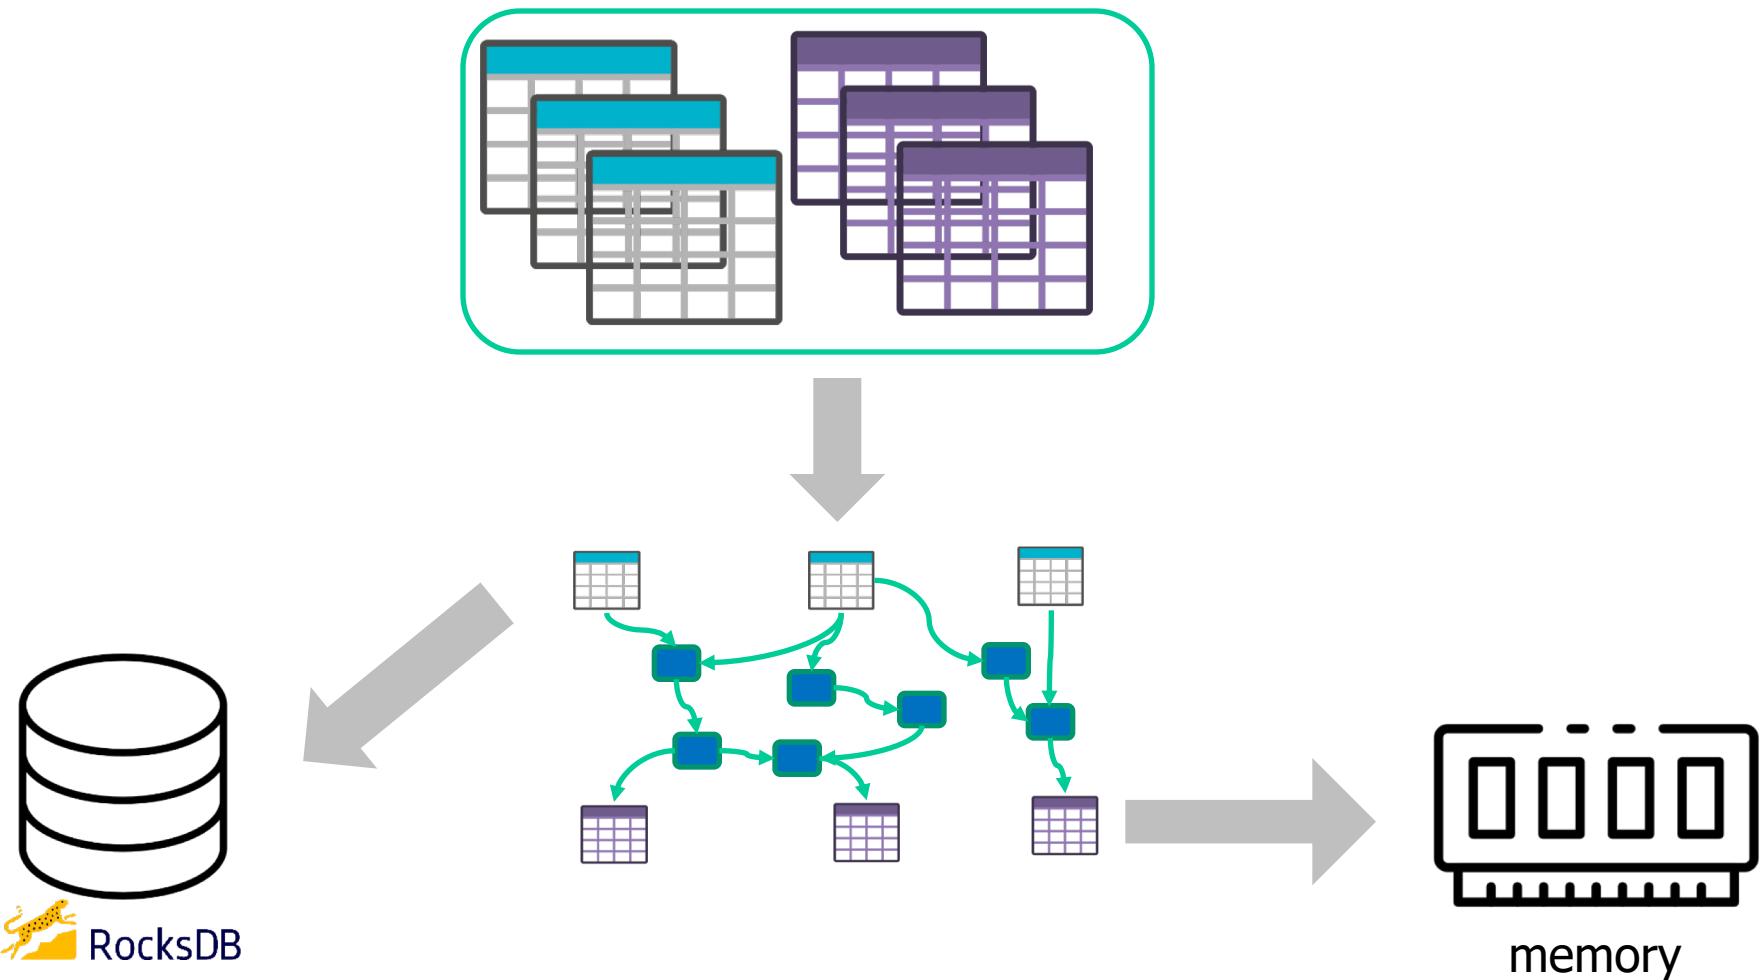


Implementation details

- Implemented in Rust
- Operation mode
 - Single server
 - Cluster of servers
- 2-way parallelism
 - Multiple shards - multiple Noria instances
 - Each instance: multiple threads to handle its shard
- Each Noria instance maintains a complete copy of the data-flow graph, but holds only the state of its operator shards
- Zookeeper: Controller election
 - Changes in data flow, shards assignment

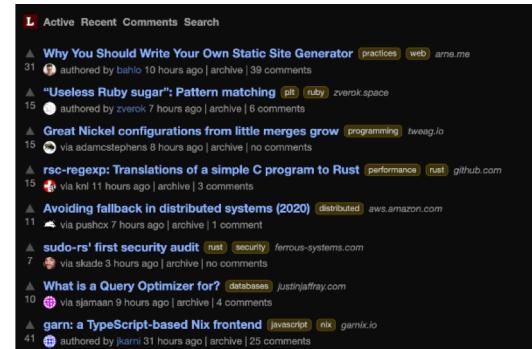


Implementation details



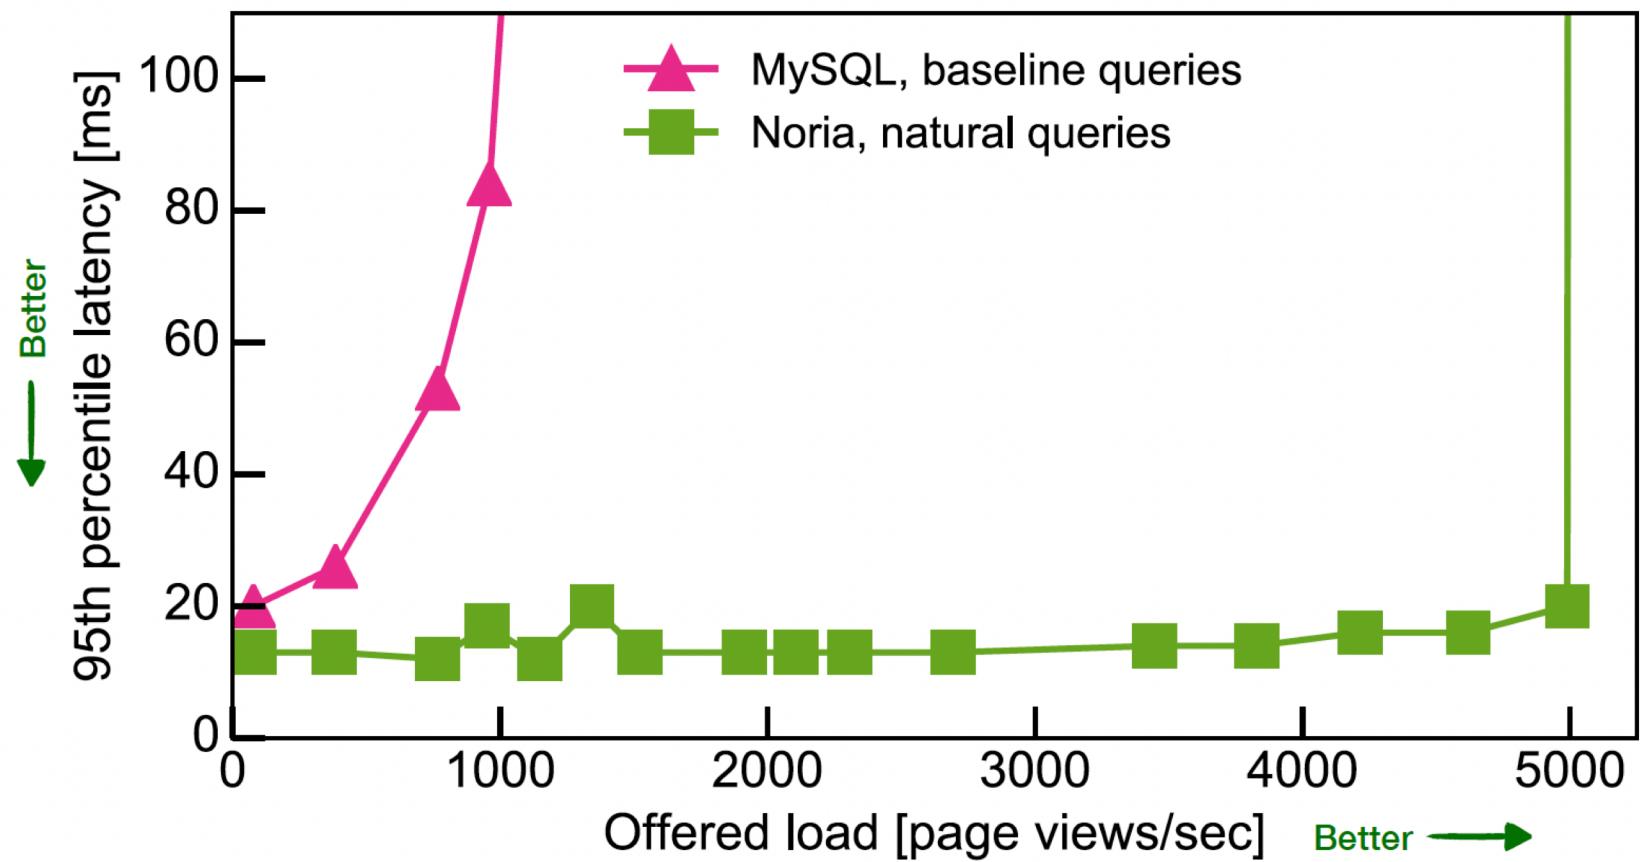
Evaluation

- Lobsters (<http://lobste.rs>)
 - 9.2K users & 40K Stories & 120K Comments
 - Emulating the production load
- Ruby-on-Rails web application with MySQL backend
- Hand-optimized SQL queries by developers to precompute various aggregations
- Noria data flow: 235 operators & 35 views
- Experiments
 - Server: Amazon EC2 c5.4xlarge instance - 16 vCPUs
 - Client: Amazon EC2 c5.4xlarge instance



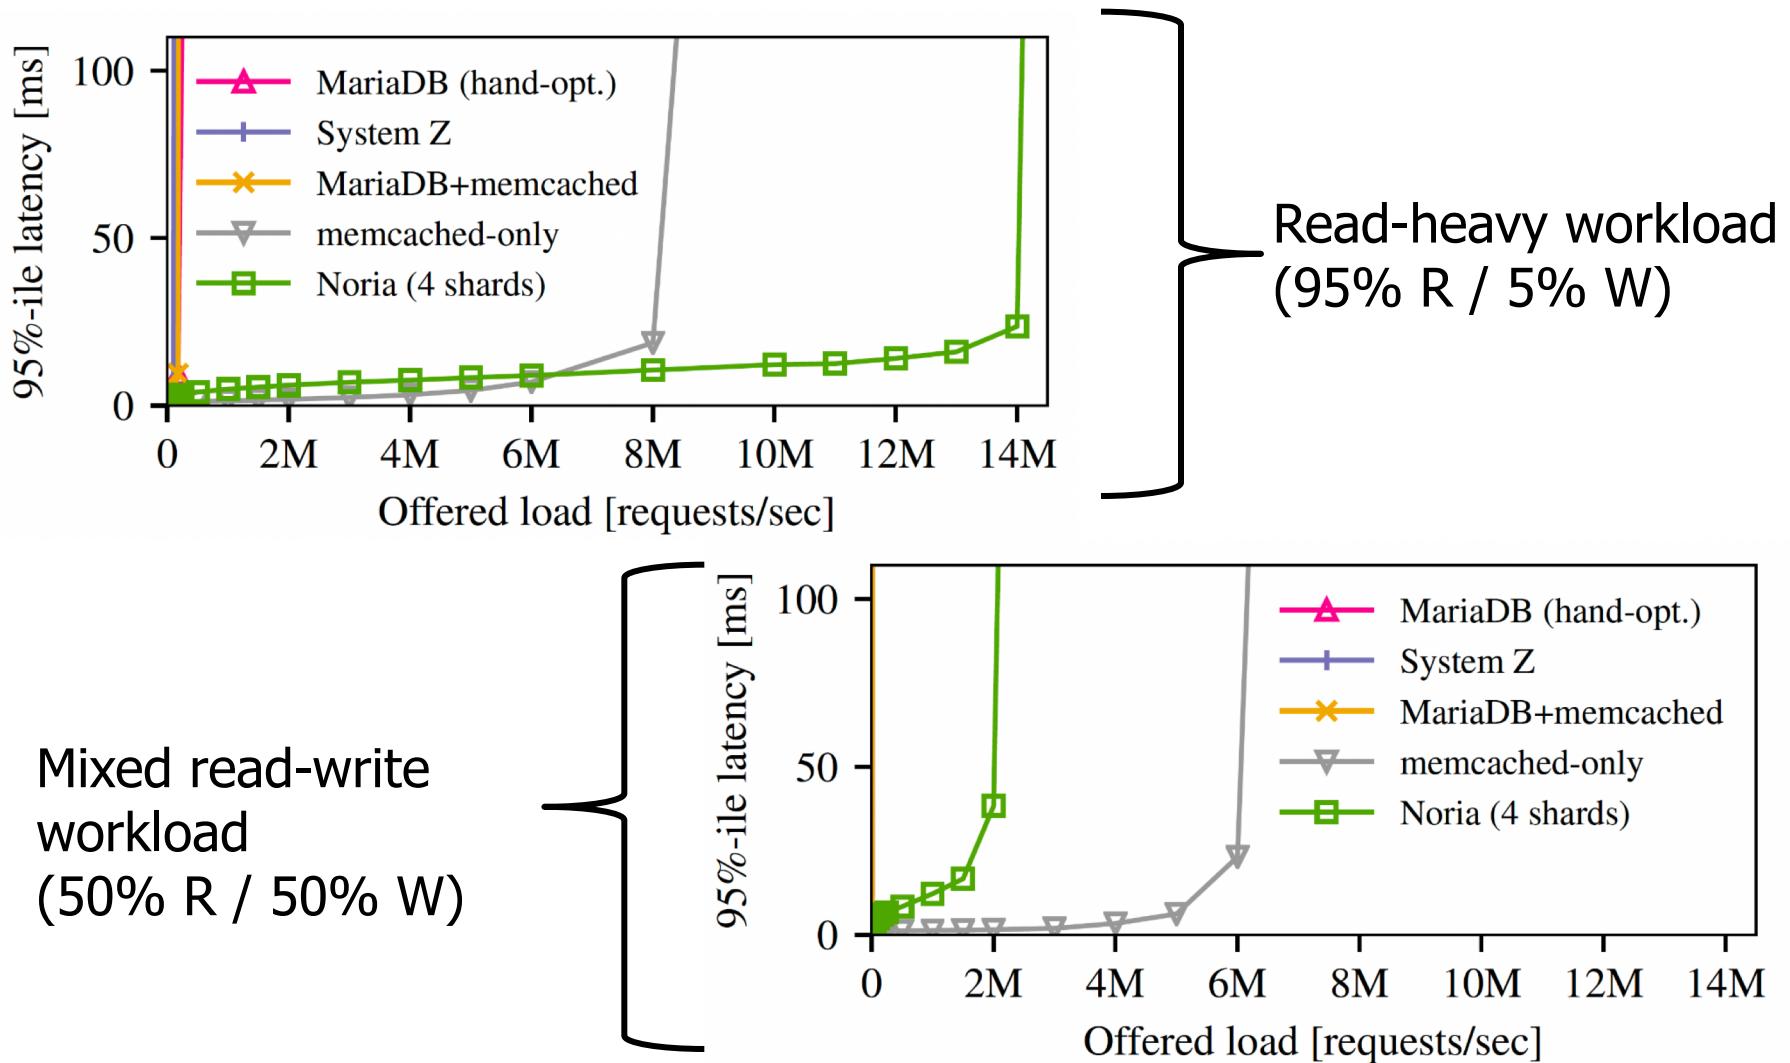
Evaluation - cont'd

- Noria performance improvements for real web applications



Evaluation - cont'd

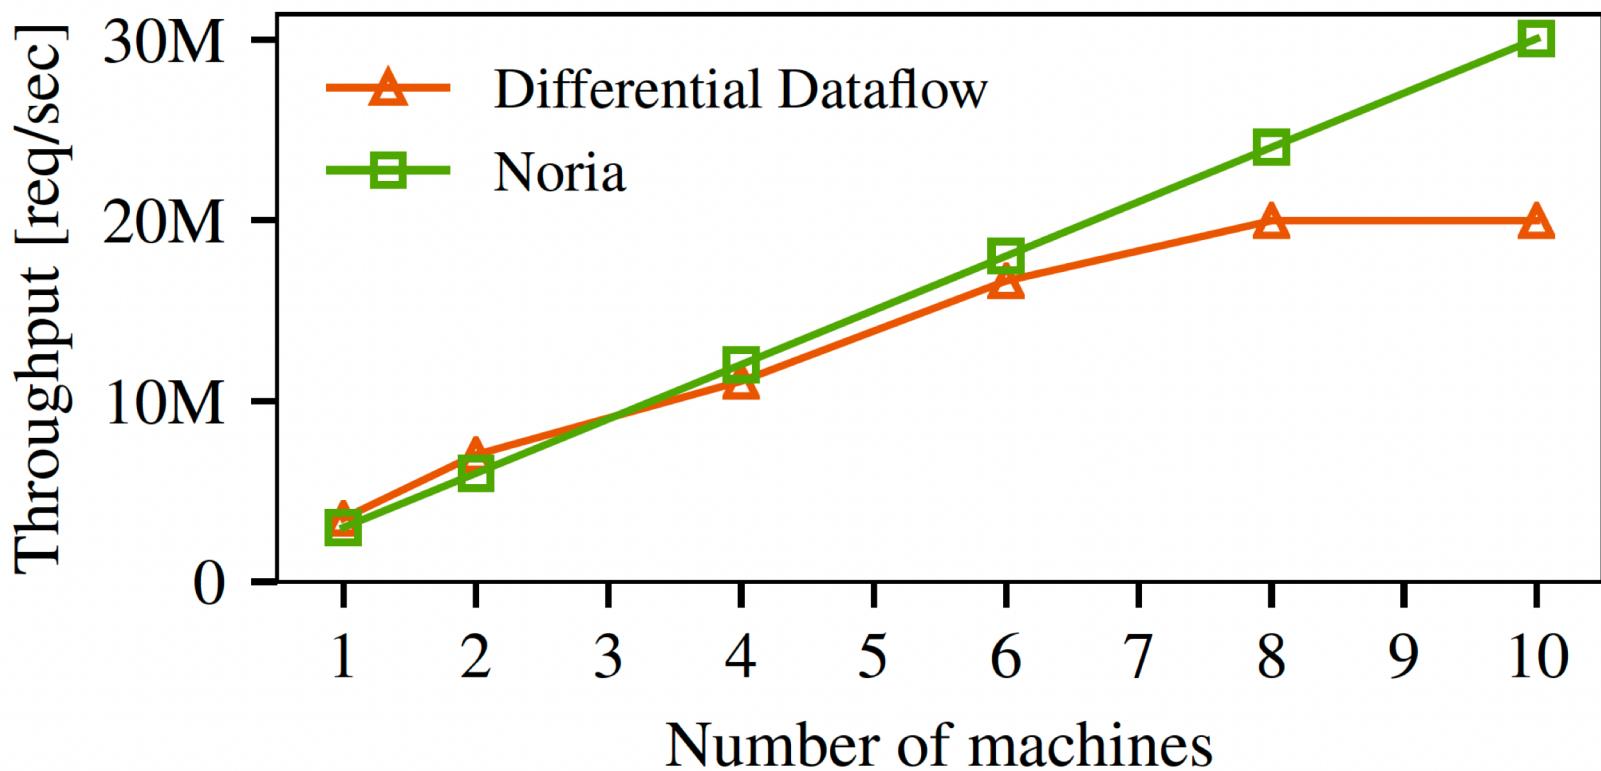
- Comparison of Noria with other alternatives



Mixed read-write workload
(50% R / 50% W)

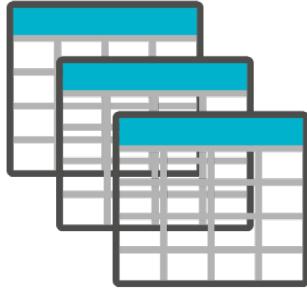
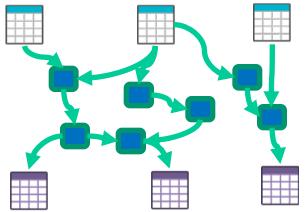
Evaluation - cont'd

- Utilization of multiple servers



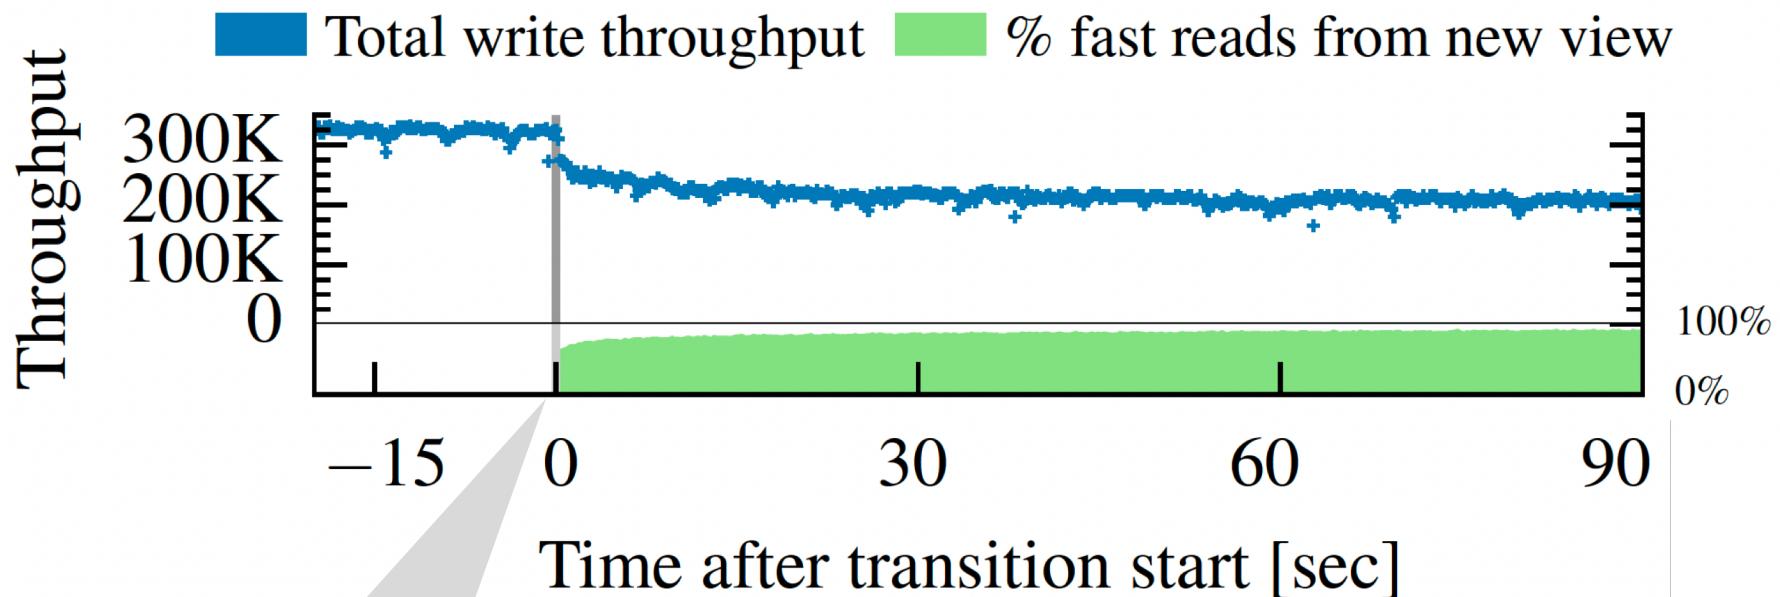
Evaluation - cont'd

- Space overhead of the data-flow and performance with limited resources
- Goal: <100ms latency at 2,300 pages/sec (Lobsters production load)

	Lobsters size	x10 Lobsters size
 base tables		
 data flow	137 MB 525 MB 60% of all state	866 MB 2.6 MB 38% of all state

Evaluation - cont'd

- Adaptability of Noria to new queries and schema updates



Conclusion

- A partially-stateful data flow model tailored for read-heavy web applications
- Fast reads with restricted memory footprint
- Near instantaneous dynamic transition for data flow graphs, in response to query changes
- Prototype implementation which is backwards-compatible with MySQL protocol interface
- With limitations
 - It only guarantees eventual consistency
 - Lacks support for some SQL keywords

References and Links

- Citation
 - Gjengset, J., Schwarzkopf, M., Behrens, J., Araújo, L.T., Ek, M., Kohler, E., Kaashoek, M.F. and Morris, R., 2018. Noria: dynamic, partially-stateful data-flow for high-performance web applications. In 13th USENIX Symposium on Operating Systems Design and Implementation (OSDI 18) (pp. 213-231).
- Link to paper and slides
 - <https://www.usenix.org/conference/osdi18/presentation/gjengset>
- Source code
 - <https://github.com/mit-pdos/noria>
- Link to these slides
 - <https://github>